Yiqiao Jin (Ahren) — Last Updated: April 18th, 2026

Contact

Georgia Institute of Technology
Atlanta, GA 30332
WeChat: go_bruins
x.com/Ahrenjin
github.com/Ahren09
(470)962-0241

ahren09.github.io
linkedin.com/in/ahren-jin
Google Scholar: eY85qm4AAAAJ
Semantic Scholar: 2087723977
ORCID: 0000-0002-6974-5970
yjin328 [AT] gatech.edu
ahren2040 [AT] g.ucla.edu

Research Interests

- Large Language Models, LLM Agents, Multimodal Learning, Graph Neural Networks, Computational Social Science.

Education

Georgia Institute of Technology — Aug. 2022 – Present

- Ph.D. Candidate, Computer Science. GPA: **4.0/4.0**.

University of California, Los Angeles (UCLA) — Sep. 2018 – Dec. 2021

- B.S., Computer Science. GPA: **3.82/4.0**.
- Published 4 papers (including 3 first-authored papers) at top-tier machine learning and data mining conferences (AAAI, KDD, Web Conference).

Research Experience

Georgia Tech College of Computing, School of CSE — Aug 2022 – Present
Graduate Research Assistant — Atlanta, GA
Advisor: Dr. Srijan Kumar

- Research Topics: Large Language Models (EMNLP'25, ACL'25, WebConf'24, ACL'24), LLM Robustness and Safety, Multimodal LLMs (ACL'24, Under Review at ACL'26), Recommender Systems and Dynamic Graph Mining (KDD'23), Social Network Analysis (CIKM'24, KDD'23), Fair Graph Mining (CIKM'24).

J.P. Morgan AI Research — May 2025 – Aug 2025
Research Scientist Intern — New York, NY

- Developed SlideAgent, an advanced framework based on Multimodal Large Language Models (MLLMs) and GUI parsing tools for understanding complex infographics. It assists human users in navigating and extracting real-time insights from slide decks and financial reports.
- Engineered LLM agentic systems for parsing high-resolution visual data, focusing on fine-grained elements like charts, tables, and text blocks.
- Integrated Context Compression & Summarization and Retrieval-augmented Generation (RAG) techniques to handle long-context efficiently under strict computational budgets.
- Authored paper *SlideAgent: Hierarchical Agentic Framework for Multi-Page Visual Document Understanding* is under review at ACL ARR.
- Mentors: Rachneet Kaur, Zhen Zeng. Managers: Sumitra Ganesh, Manuela Veloso.

Visa Research — Aug 2024 – May 2025
Research Collaborator (Part-time) — Remote

- Developed SARA, a unified Retrieval-augmented Generation (RAG) framework to balance local factual precision with global knowledge coverage under strict context budgets.
- Proposed a novel hybrid compression strategy that represents external evidence through both fine-grained natural-language spans and compact, interpretable semantic compression vectors.
- Engineered an iterative context refinement mechanism that utilizes compression vectors for dynamic reranking, effectively reducing document redundancy while maximizing query informativeness.
- Achieved significant performance gains across multiple datasets and open-source LLM families (Mistral, Llama, Gemma).

- Authored paper *SARA: Selective and Adaptive Retrieval-augmented Generation with Context Compression* is under review at ACL ARR.
- Mentors: Vineeth Rakesh Mohan, Yingtong Dou, Menghai Pan. Manager: Mahashweta Das.

Adobe Research May 2024 – Aug 2024
Research Scientist Intern San Jose, CA

- Developed ScreenLLM, a specialized multimodal LLM for Graphical User Interface (GUI) understanding and action prediction.
- Proposed a novel stateful screen schema that represents dynamic user sessions as compact, time-aware textual summaries.
- Engineered a high-efficiency key frame extraction method based on second-order pixel changes to isolate significant UI transitions like pop-up menus.
- Published 1 paper *ScreenLLM: Stateful Screen Schema for Efficient Action Understanding and Prediction* at the Web Conference 2025 (WWW'25) MM4SG workshop.
- Mentors: Gang Wu, Yu Shen, Stefano Petrangeli. Managers: Saayan Mitra, Vishy Swaminathan

Microsoft Research Asia (MSRA), Social Computing Group Dec. 2020 – July 2022
Research Scientist Intern Beijing, China

- Research Topics: Large Language Models (EMNLP'24, ICML'24, ICML'23, AAAI'23), LLM Agents (EMNLP'24, ICML'24), Multicultural and Multimodal LLM, Scientometric Analysis, Computational Social Science, Misinformation Detection (KDD'22, AAAI'22), Few-shot Learning (ACL'24, AAAI'23) , Explainable AI (AAAI'22).
- Mentors: Xiting Wang, Jindong Wang, and Xing Xie.

UCLA, Scalable Analytics Institute (ScAi) June 2021 – June 2022
Undergraduate Research Assistant Los Angeles, CA

- Research Topics: Large Language Models (EMNLP'25, ACL'25, EMNLP'24), Graph Neural Networks and Data Mining (WWW'23), LLM Fine-tuning, Recommender Systems (WWW'23).
- I am continuously collaborating with UCLA researchers during my Ph.D. studies.
- Advisors: Yizhou Sun and Wei Wang.

Publications

Conference Papers (First-Authored)

1. **Yiqiao Jin**, Rachneet Kaur, Zhen Zeng, Sumitra Ganesh, and Srijan Kumar. *SlideAgent: Hierarchical Agentic Framework for Multi-Page Visual Document Understanding.* In Proceedings of ACL'26 main conference.
   Acceptance rate: 19.0%.
2. **Yiqiao Jin**, Kartik Sharma, Vineeth Rakesh, Yingtong Dou, Menghai Pan, Mahashweta Das, and Srijan Kumar. *SARA: Selective and Adaptive Retrieval-augmented Generation with Context Compression.* In Proceedings of ACL'26 main conference.
   Acceptance rate: 19.0%.
3. **Yiqiao Jin***, Qinlin Zhao*, Yiyang Wang, Hao Chen, Kaijie Zhu, Yijia Xiao, Jindong Wang. *AgentReview: Exploring Peer Review Dynamics with LLM Agents.* In Proceedings of EMNLP'24 main conference **Oral Presentation**. Presented at SoCal NLP Symposium 2024.
   Acceptance rate: 20.8%.
4. **Yiqiao Jin**, Minje Choi, Gaurav Verma, Jindong Wang, Srijan Kumar. *MM-Soc: Benchmarking Multimodal Large Language Models in Social Media Platforms.* In Proceedings of ACL'24 and the 3rd ALVR Workshop.
   Acceptance rate: 23.5%.
5. **Yiqiao Jin***, Mohit Chandra*, Gaurav Verma, Yibo Hu, Munmun De Choudhury, Srijan Kumar. *Better to Ask in English: Cross-Lingual Evaluation of Large Language Models for Healthcare Queries.* In Proceedings of the ACM Web Conference 2024. **Oral Presentation**.
   Acceptance rate: 20.2%.
   Media Coverage by Scientific American, The World, and Georgia Tech.

6. **Yiqiao Jin**, Yeon-Chang Lee, Kartik Sharma, Meng Ye, Karan Sikka, Ajay Divakaran, Srijan Kumar. *Predicting Information Pathways Across Online Communities.* In Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD'23). **Oral Presentation**.
   Acceptance rate: 22.1%.

7. **Yiqiao Jin**, Yunsheng Bai, Yanqiao Zhu, Yizhou Sun, Wei Wang. *Code Recommendation for Open Source Project Developers.* In Proceedings of the ACM Web Conference 2023. **Oral Presentation**.
   Acceptance rate: 19.2%

8. **Yiqiao Jin**, Xiting Wang, Yaru Hao, Yizhou Sun, Xing Xie. *Prototypical Fine-tuning: Towards Robust Performance Under Varying Data Sizes.* In Proceedings of the 37th AAAI Conference (AAAI'23). **Oral Presentation**.
   Acceptance rate: 19.6%.

9. **Yiqiao Jin**, Xiting Wang, Ruichao Yang, Yizhou Sun, Wei Wang, Hao Liao, Xing Xie. *Towards Fine-Grained Reasoning for Fake News Detection.* In Proceedings of the 36th AAAI Conference (AAAI'22). **Oral Presentation**.
   Acceptance rate: 14.6%.

Conference Papers (Collaborative Projects)

1. Mohit Chandra, Siddharth Sriraman, Harneet Singh Khanuja, **Yiqiao Jin**, and Munmun De Choudhury. *Reasoning Is Not All You Need: Examining LLMs for Multi-Turn Mental Health Conversations*. In Proceedings of ACL'26 main conference.
   Acceptance rate: 19.0%.

2. Hanoz Bhathena, Parin Rajesh Jhaveri, Rohan Mittal, Prateek Singh, Aymen Kallala, Rachneet Kaur, **Yiqiao Jin**, Zhen Zeng, Adwait Ratnaparkhi, and Denis Kochedykov. *MM-BizRAG: Rethinking Multimodal Retrieval-Augmented Generation for General Purpose Enterprise Q&A*. In Proceedings of ACL'26 Industry Track.

3. Kartik Sharma, **Yiqiao Jin**, Vineeth Rakesh, Yingtong Dou, Menghai Pan, Mahashweta Das, and Srijan Kumar. *Sysformer: Safeguarding Frozen Large Language Models with Adaptive System Prompts.* In Proceedings of the 14th International Conference on Learning Representations (ICLR'26).
   Acceptance rate: 28%.

4. Guancheng Wan*, Lucheng Fu*, Haoxin Liu, **Yiqiao Jin**, Hejia Geng, Eric Hanchen Jiang, Hui Yi Leong, Jinhe Bi, Yunpu Ma, Xiangru Tang, B. Aditya Prakash, Yizhou Sun, Wei Wang. *Beyond Magic Words: Sharpness-Aware Prompt Evolving for Robust Large Language Models.* In Proceedings of the 14th International Conference on Learning Representations (ICLR'26).
   Acceptance rate: 28%.

5. Vibhor Agarwal, **Yiqiao Jin**, Mohit Chandra, Munmun De Choudhury, Srijan Kumar, Nishanth Sastry *MedHalu: Hallucinations in Responses to Healthcare Queries by Large Language Models.* In Proceedings of ICWSM'26.

6. Yijia Xiao, Wanjia Zhao, Junkai Zhang, **Yiqiao Jin**, Han Zhang, Zhicheng Ren, Renliang Sun, Haixin Wang, Guancheng Wan, Pan Lu, Xiao Luo, Yu Zhang, James Zou, Yizhou Sun, Wei Wang. *Protein Large Language Models: A Comprehensive Survey*. In Proceedings of EMNLP'25.
   Acceptance rate: 22.2%.

7. Junyu Luo, Bohan Wu, Xiao Luo, Zhiping Xiao, **Yiqiao Jin**, Rong-Cheng Tu, Nan Yin, Yifan Wang, Jingyang Yuan, Wei Ju, Ming Zhang. *A Survey on Efficient LLM Training: From Data-centric Perspectives.* In Proceedings of ACL'25 main conference.

8. Qinlin Zhao, Jindong Wang, Yixuan Zhang, **Yiqiao Jin**, Kaijie Zhu, Hao Chen, Xing Xie. *CompeteAI: Understanding the Competition Behaviors in Large Language Model-based Agents.* In Proceedings of the 41st International Conference on Machine Learning (ICML'24), **Oral Presentation**.
   Acceptance rate: 27.5%.

9. Yijia Xiao, **Yiqiao Jin**, Yushi Bai, Yue Wu, Xianjun Yang, Xiao Luo, Wenchao Yu, Xujiang Zhao, Yanchi Liu, Haifeng Chen, Wei Wang, Wei Cheng. *Large Language Models Can Be Good Privacy Protection Learners.* In Proceedings of EMNLP'24 main conference and SoCal NLP

Symposium 2024.
Acceptance rate: 20.8%.

10. Jinghan Zhang, Xiting Wang, **Yiqiao Jin**, Changyu Chen, Xinhao Zhang, Kunpeng Liu. *Prototypical Reward Network for Data-Efficient RLHF.* In Proceedings of ACL'24. **Oral Presentation**.
Acceptance rate: 23.5%.

11. Neng Kai Nigel Neo*, Yeon-Chang Lee*, **Yiqiao Jin**, Sang-Wook Kim, Srijan Kumar. *Towards Fair Graph Anomaly Detection: Problem, New Datasets, and Evaluation.* In Proceedings of the 33rd ACM International Conference on Information and Knowledge Management (CIKM'24).
Acceptance rate: 23.2%.

12. Changyu Chen, Xiting Wang, **Yiqiao Jin**, Victor Ye Dong, Li Dong, Rui Yan, Jim Cao, Yi Liu. *Semi-Offline Reinforcement Learning for Optimized Text Generation.* In Proceedings of the 40th International Conference on Machine Learning (ICML'23).
Acceptance rate: 27.9%.

13. Ruichao Yang, Xiting Wang, **Yiqiao Jin**, Chaozhuo Li, Jianxun Lian, Xing Xie. *Reinforcement Subgraph Reasoning for Fake News Detection.* In Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD'22).
Acceptance rate: 14.9%.

Journals

1. Chengyuan Deng, Yiqun Duan, Xin Jin, Heng Chang, Yijun Tian, Han Liu, Henry Peng Zou, **Yiqiao Jin**, Yijia Xiao, Yichen Wang, Shenghao Wu, Zongxing Xie, Kuofeng Gao, Sihong He, Jun Zhuang, Lu Cheng, Haohan Wang. *Deconstructing The Ethics of Large Language Models from Long-standing Issues to New-emerging Dilemmas.* Accepted at AI and Ethics.

Workshops

1. **Yiqiao Jin**, Yijia Xiao, Yiyang Wang, and Jindong Wang. *SciEvo: A 2 Million, 30-Year Cross-disciplinary Dataset for Temporal Scientometric Analysis.* Best Paper Award (2 out of 25 accepted papers) at AAAI'25 Good-Data Workshop.

2. **Yiqiao Jin**, Andrew Zhao, Yeon-Chang Lee, Meng Ye, Ajay Divakaran, and Srijan Kumar. *Empowering Interdisciplinary Insights with Dynamic Graph Embedding Trajectories.* Accepted at KDD'25 TGL Workshop.

3. **Yiqiao Jin**, Gang Wu, Yu Shen, and Stefano Petrangeli, *Stateful Screen Schema for Efficient Action Understanding and Prediction.* Accepted at WebConf'25 MM4SG Workshop.

4. Yijia Xiao, Edward Sun, **Yiqiao Jin**, Qifan Wang, and Wei Wang. *ProteinGPT: Multimodal LLM for Protein Property Prediction and Structure Understanding.* Accepted at ICML'25 FM4LS Workshop.

5. Shudong Liu, **Yiqiao Jin**, Cheng Li, Derek F. Wong, Qingsong Wen, Lichao Sun, Haipeng Chen, Xing Xie, and Jindong Wang. *CultureVLM: Characterizing and Improving Cultural Understanding of Vision-Language Models for over 100 Countries.* Accepted at CVPR'25 VLMs4All Workshop.

6. Sejoon Oh*, **Yiqiao Jin***, Megha Sharma, Donghyun Kim, Gaurav Verma, Eric Ma, Srijan Kumar. *UniGuard: Multimodal Safety Guardrails for Jailbreak Attacks on Multimodal Large Language Models.* Accepted at AAAI'25 DAI Workshop.

7. Yijia Xiao, Edward Sun, **Yiqiao Jin**, Wei Wang. *RNA-GPT: Multimodal Generative System for RNA Sequence Understanding.* Accepted at NeurIPS'24 MLSB Workshop.

Preprints

1. Yiyang Wang, **Yiqiao Jin**, Alex Cabral, and Josiah Hester. *MASCOT: Towards Multi-Agent Socio-Collaborative Companion Systems*. Under Review at ACL ARR.

2. Yinyi Luo, **Yiqiao Jin**, Weichen Yu, Mengqi Zhang, Srijan Kumar, Xiaoxiao Li, Weijie Xu, Xin Chen, Jindong Wang. *AgentArk: Distilling Multi-Agent Intelligence into a Single LLM Agent*. Under Review at ICML'26.

3. Zhaolong Su*, Yinyi Luo*, **Yiqiao Jin***, Mengqi Zhang, Wenyue Hua, Srijan Kumar, Qingsong Wen, Jindong Wang *Consistency Should Be the Priority for Unified Multimodal Models*. Under Review at ICML'26.

4. Jiaxi Yang*, Mengqi Zhang*, **Yiqiao Jin***, Hao Chen, Qingsong Wen, Lu Lin, Yi He, Srijan Kumar, Weijie Xu, James Evans, and Jindong Wang. *Topological Structure Learning Should Be A Research Priority for LLM-Based Multi-Agent Systems*. Under Review at COLM'26.
5. Kartik Sharma, **Yiqiao Jin**, Rakshit Trivedi, Srijan Kumar. *Efficient Knowledge Probing of Large Language Models by Adapting Pre-trained Embeddings.* Under Review at AAAI'26.
6. Ye Yuan, Junyu Luo, Guancheng Wan, Jinsheng Huang, Chengwu Liu, Junwei Yang, Yifang Qin, Zhiping Xiao, QingqingLong, Meng Xiao, **Yiqiao Jin**, Jianhong Tu, Yuqi Chen, Wei Ju, Zhongwei Wan, Yusheng Zhao, Xiao Luo, Yiwei Fu, Yizhou Sun, Wei Wang, Chenguang Wang, and Ming Zhang. *A Comprehensive Survey of Evaluating Multimodal Foundation Models: Hierarchical Perspective and Extensive Applications* Under Review at ACL ARR.

Professional Experience

Amazon, Fulfillment By Amazon (FBA) — June 2020 – Sep. 2020
Software Engineer Intern — Seattle, WA

- Designed and implemented `IAR Manual Analysis`, a scalable and efficient workflow using `AWS Step Functions` and `AWS Lambda`. This service automates the aggregation of data points from multiple sources like `Amazon S3` and `DynamoDB` for `SageMaker` ML model training, handling ≥ 16,000 requests per summary stage;
- Automated the deployment of the workflow across all AWS Realms (EU/FE/NA) through `CloudFormation`;
- Establish `DataCraft` pipeline to enable automatic data ingestion from `DynamoDB` into the `Andes` dataset catalog, promoting broader internal adoption of these datasets for cross-functional teams and enhancing data accessibility;
- Perform `ablation analysis` on the inventory reconciliation model, identifying key performance bottlenecks and optimizing model performance.

IBM, China Development Laboratories — June 2019 – Sep. 2019
Software Engineer Intern — Beijing, China

- Developed `Compass DataRouter`, a routing service for `Compass` project based on `Golang` and `MongoDB`, reducing memory usage and accelerating data retrieval;
- Enhanced the monitoring dashboard for `Compass` towards a more intuitive and responsive user interface with `React.js`.

Services

Georgia Tech School of Computer Science Graduate Student Association (SCSGSA) — Aug. 2024 –
Communications Chair — Atlanta, GA

- Co-organize the College of Computing (CoC) Graduate Welcome Event with over 1,000 attendees, including new students, faculty, and alumni;
- Lead the design and maintenance of the SCSGSA website (`https://scsgsa.cc.gatech.edu/`);
- Promote SCSGSA events, including workshops, networking sessions, and professional development programs;

Academic Services

- Area Chair: ICLR'24/23 Tiny Paper Track
- PC Member/Reviewer
  - **Conferences**:
    - AAAI 26/25/24/23;
    - ACL ARR (Since Feb. 2024. Including ACL/EMNLP/NAACL);
    - AISTATS 25/24;
    - ASONAM 24/23;
    - CIKM 24/23;
    - COLM 26/25;
    - CVPR 26/25;
    - ICCV 25;
    - ICLR 26/25/24, ICLR 24/23 Tiny Paper;
    - ICML 26/25/24;

    * KDD 25/24/23;
    * LoG 24/23/22;
    * NeurIPS 26/25/24/23, NeurIPS 25/24/23 Datasets & Benchmarks Track;
    * SDM 24;
    * Web Conference 25/24.
  - **Journals**:
    * ACM Transactions on Information Systems (TOIS); 2023 –
    * ACM Transactions on Knowledge Discovery from Data (TKDD); 2023 –
    * IEEE Transactions on Knowledge and Data Engineering (TKDE); 2022 –
    * ACM Transactions on Recommender Systems (TORS); 2022 –
    * ACM Transactions on Intelligent Systems and Technology (TIST); 2022 –
    * ACM Transactions on Social Computing (TSC); 2022 –
    * IEEE Transactions on Artificial Intelligence (TAI); 2023 –
    * PLOS ONE; 2023 –
    * International Journal of Data Science and Analytics (JDSA); 2021 –
    * SCIENCE CHINA Information Sciences (SCIS); 2022 –
    * The Journal of Data-centric Machine Learning Research (DMLR). 2023 –
  - **Workshops**
    * AFT@NeurIPS2023, Ai4Science@NeurIPS2023, Ai4D3@NeurIPS2023, GenBio@NeurIPS2023, TGL Workshop@NeurIPS2023;
    * SPIGM@ICML2023.

- Volunteer: NAACL'25, AAAI'25, SC'24, VLDB'24

Honors and Awards

- Best Paper Award at Good-Data @ AAAI'25 Workshop. 2025
- Roblox Graduate Fellowship Finalist. 2024
- AAAI Student Scholarship (3 times). 2025, 2023, 2022
- Microsoft Research "Star of Tomorrow" Award of Excellence. 2021
- UCLA Dean's Honor List for Superior Academic Achievement (5 times). 2019 – 2021
  - Spring 2019, Winter 2020, Spring 2020, Winter 2021, Spring 2021