Yiqiao Jin (Ahren) **Last Updated: Aug. 24th, 2026**

Contact

Georgia Institute of Technology, Atlanta, GA 30332
Website: ahren09.github.io
Google Scholar: scholar.google.com/citations?user=eY85qm4AAAAJ
Semantic Scholar: 2087723977
X: x.com/Ahrenjin
Email: yjin328@gatech.edu
LinkedIn: linkedin.com/in/ahren-jin
GitHub: github.com/Ahren09
ORCID: 0000-0002-6974-5970
Phone: (470)962-0241
Email: ahren2040@g.ucla.edu

Research Interests

- Large Language Models (LLMs); Multimodal Large Language Models (MLLMs) and Vision-Language Models (VLMs); Post-Training and Reinforcement Learning; LLM Agents and Multi-Agent Systems; Multimodal Reasoning, Interaction, and Evaluation.

Education

Georgia Institute of Technology Aug. 2022 – Present

- Ph.D. Candidate, Computer Science. GPA: **4.0/4.0**.
- Selected Publications: **MASCOT** (EMNLP 2026 main conference), **CultureVLM** (EMNLP 2026), **SARA** (ACL 2026 Oral), **SlideAgent** (conference), **MedAgent** (ACL 2026 main conference), **Sysformer** (ICLR 2026), **TARE** (ICLR 2026), **MedHalu** (ICWSM 2026), **PrivacyMind** (EMNLP 2024 main conference), **AgentReview** (EMNLP 2024 Oral), **ProtoRM** (ACL 2024 main conference), **MM-Soc** (ACL 2024), **XLingEval** (Web Conference 2024 Oral), FairGAD (CIKM 2024), and **INPAC** (KDD 2023 Oral).

University of California, Los Angeles (UCLA) Sep. 2018 – Dec. 2021

- B.S., Computer Science. GPA: **3.82/4.0**.
- Selected Publications: **CODER** (Web Conference 2023 Oral), **FinerFact** (AAAI 2022 Oral), **SureFact** (KDD 2022 Oral), **Prototypical Fine-tuning** (AAAI 2023 Oral).

Research Experience

Georgia Tech College of Computing, School of CSE Aug. 2022 – Present
Graduate Research Assistant Atlanta, GA
Advisor: Dr. Srijan Kumar

- Research Topics: Large Language Models (ACL2026, ICWSM2026, EMNLP2025, ACL2025, WebConf2024, ACL2024), LLM Robustness and Safety, Multimodal LLMs (ACL2024), Recommender Systems and Dynamic Graph Mining (KDD2023), Social Network Analysis (CIKM2024, KDD2023), Fair Graph Mining (CIKM2024).

Amazon Aug. 2026 – Nov. 2026
Applied Scientist Intern Santa Clara, CA

- Develop a unified agentic framework based on diffusion language models (dLMs) for task completion and world modeling.
- Mentor: Anwesan Pal. Manager: Narayanan Sadagopan .

Waymo May 2026 – Aug. 2026
Ph.D. Intern Mountain View, CA

- Developed a unified VLM for autonomous driving perception (**patent application pending**) that learns dense metric scene geometry and sensor-grounded spatial reasoning from camera observations and sparse LiDAR. Patent application is in-progress.
- Built core perceptual capabilities for vision–language–action (VLA) systems—dense depth, object-centric distance, driving relevance, depth ordering, and time-to-reach reasoning—through a shared geometry–language scene representation.
- Designed a three-stage training recipe that anchors metric scale on 160K Waymo frames, converts sparse LiDAR ($< 1.1\%$ pixel coverage) into reliability-weighted dense supervision, and instruction-tunes the VLM while preserving geometric perception.
- Curated 14K spatial instruction examples grounded in LiDAR and ego-motion, and studied temporal (± 0.1 s) and multi-camera context to diagnose causal token ordering and viewpoint shifts in video-based scene representations.
- Mentors: Mayank Singal, Prasanna Krishnasamy. Managers: Ming Zou, Christian Lauterbach.

J.P. Morgan AI Research June 2025 – Aug 2025
Research Scientist Intern New York, NY

- Developed SlideAgent, an advanced framework based on Multimodal Large Language Models (MLLMs) and GUI parsing tools for understanding complex infographics (**patent application pending**). It assists human users in navigating and extracting real-time insights from slide decks and financial reports. Patent application is in-progress.

- Engineered LLM agentic systems for parsing high-resolution visual data, focusing on fine-grained elements like charts, tables, and text blocks.
- Integrated Context Compression & Summarization and Retrieval-augmented Generation (RAG) techniques to handle long-context efficiently under strict computational budgets.
- First-authored a paper *SlideAgent: Hierarchical Agentic Framework for Multi-Page Visual Document Understanding* accepted at ACL2026 main conference.
- Mentors: Rachneet Kaur, Zhen Zeng. Managers: Sumitra Ganesh, Manuela Veloso.

Visa Research Aug. 2024 – May 2025
Research Collaborator (Under Visa Research Grants) Remote

- Developed SARA, a unified Retrieval-augmented Generation (RAG) framework to balance local factual precision with global knowledge coverage under strict context budgets.
- Proposed a novel hybrid compression strategy that represents external evidence through both fine-grained natural-language spans and compact, interpretable semantic compression vectors.
- Engineered an iterative context refinement mechanism that utilizes compression vectors for dynamic reranking, effectively reducing document redundancy while maximizing query informativeness.
- Achieved significant performance gains across multiple datasets and open-source LLM families (Mistral, Llama, Gemma).
- First-authored a paper *SARA: Selective and Adaptive Retrieval-augmented Generation with Context Compression* accepted at ACL 2026 main track (oral presentation).
- Mentors: Vineeth Rakesh Mohan, Yingtong Dou, Menghai Pan. Manager: Mahashweta Das.

Adobe Research May 2024 – Aug 2024
Research Scientist Intern San Jose, CA

- Developed ScreenLLM, a specialized multimodal LLM for Graphical User Interface (GUI) understanding and action prediction.
- Proposed a novel stateful screen schema that represents dynamic user sessions as compact, time-aware textual summaries.
- Engineered a high-efficiency key frame extraction method based on second-order pixel changes to isolate significant UI transitions like pop-up menus.
- Published 1 paper *ScreenLLM: Stateful Screen Schema for Efficient Action Understanding and Prediction* at WebConf2025 MM4SG Workshop.
- Mentors: Gang Wu, Yu Shen, Stefano Petrangeli. Managers: Saayan Mitra, Vishy Swaminathan

Microsoft Research Asia (MSRA), Social Computing Group Dec. 2020 – July 2022
Research Scientist Intern Beijing, China

- Research Topics: Large Language Models (EMNLP2024, ICML2024, ICML2023, AAAI2023), LLM Agents (EMNLP2024, ICML2024), Multicultural and Multimodal LLM, Scientometric Analysis, Computational Social Science, Misinformation Detection (KDD2022, AAAI2022), Few-shot Learning (ACL2024, AAAI2023) , Explainable AI (AAAI2022).
- Mentors: Xiting Wang, Jindong Wang, and Xing Xie.

UCLA, Scalable Analytics Institute (ScAi) June 2021 – June 2022
Undergraduate Research Assistant Los Angeles, CA

- Research Topics: Large Language Models (EMNLP2025, ACL2025, EMNLP2024), Graph Neural Networks and Data Mining (WebConf2023), LLM Fine-tuning, Recommender Systems (WebConf2023).
- I am continuously collaborating with UCLA researchers during my Ph.D. studies.
- Advisors: Yizhou Sun and Wei Wang.

Publications

Conference Papers (First-Authored)

1. **Yiqiao Jin**, Rachneet Kaur, Zhen Zeng, Sumitra Ganesh, and Srijan Kumar. *SlideAgent: Hierarchical Agentic Framework for Multi-Page Visual Document Understanding.* In Proceedings of ACL2026 main conference.
Acceptance rate: 19.0%.

2. **Yiqiao Jin**, Kartik Sharma, Vineeth Rakesh, Yingtong Dou, Menghai Pan, Mahashweta Das, and Srijan Kumar. *SARA: Selective and Adaptive Retrieval-augmented Generation with Context Compression.* In Proceedings of ACL2026 main conference.
Acceptance rate: 19.0%.

3. **Yiqiao Jin***, Qinlin Zhao*, Yiyang Wang, Hao Chen, Kaijie Zhu, Yijia Xiao, Jindong Wang. *AgentReview: Exploring Peer Review Dynamics with LLM Agents.* In Proceedings of EMNLP2024 main track. **Oral Presentation**.
Presented at SoCal NLP Symposium 2024.
Acceptance rate: 20.8%.

4. **Yiqiao Jin**, Minje Choi, Gaurav Verma, Jindong Wang, Srijan Kumar. *MM-Soc: Benchmarking Multimodal Large Language Models in Social Media Platforms.* In Proceedings of ACL2024 and the 3rd ALVR Workshop.
Acceptance rate: 23.5%.

5. **Yiqiao Jin***, Mohit Chandra*, Gaurav Verma, Yibo Hu, Munmun De Choudhury, Srijan Kumar. *Better to Ask in English: Cross-Lingual Evaluation of Large Language Models for Healthcare Queries.* In Proceedings of WebConf2024. **Oral Presentation**.
Acceptance rate: 20.2%.
Media Coverage by Scientific American, The World, and Georgia Tech.

6. **Yiqiao Jin**, Yeon-Chang Lee, Kartik Sharma, Meng Ye, Karan Sikka, Ajay Divakaran, Srijan Kumar. *Predicting Information Pathways Across Online Communities.* In Proceedings of KDD2023. **Oral Presentation**.
Acceptance rate: 22.1%.

7. **Yiqiao Jin**, Yunsheng Bai, Yanqiao Zhu, Yizhou Sun, Wei Wang. *Code Recommendation for Open Source Project Developers.* In Proceedings of WebConf2023. **Oral Presentation**.
Acceptance rate: 19.2%

8. **Yiqiao Jin**, Xiting Wang, Yaru Hao, Yizhou Sun, Xing Xie. *Prototypical Fine-tuning: Towards Robust Performance Under Varying Data Sizes.* In Proceedings of AAAI2023. **Oral Presentation**.
Acceptance rate: 19.6%.

9. **Yiqiao Jin**, Xiting Wang, Ruichao Yang, Yizhou Sun, Wei Wang, Hao Liao, Xing Xie. *Towards Fine-Grained Reasoning for Fake News Detection.* In Proceedings of AAAI2022. **Oral Presentation**.
Acceptance rate: 14.6%.

Conference Papers (Collaborative Projects)

1. Yiyang Wang, **Yiqiao Jin**, Alex Cabral, and Josiah Hester. *MASCOT: Towards Multi-Agent Socio-Collaborative Companion Systems*. In Proceedings of EMNLP 2026 Main Conference.

2. Shudong Liu, **Yiqiao Jin**, Cheng Li, Derek F. Wong, Qingsong Wen, Lichao Sun, Haipeng Chen, Xing Xie, and Jindong Wang. *CultureVLM: Characterizing and Improving Cultural Understanding of Vision-Language Models for over 100 Countries.* In Proceedings of EMNLP 2026.

3. Mohit Chandra, Siddharth Sriraman, Harneet Singh Khanuja, **Yiqiao Jin**, and Munmun De Choudhury. *Reasoning Is Not All You Need: Examining LLMs for Multi-Turn Mental Health Conversations*. In Proceedings of ACL2026 main conference.
Acceptance rate: 19.0%.

4. Hanoz Bhathena, Parin Rajesh Jhaveri, Rohan Mittal, Prateek Singh, Aymen Kallala, Rachneet Kaur, **Yiqiao Jin**, Zhen Zeng, Adwait Ratnaparkhi, and Denis Kochedykov. *MM-BizRAG: Rethinking Multimodal Retrieval-Augmented Generation for General Purpose Enterprise Q&A*. In Proceedings of ACL2026 Industry Track.

5. Kartik Sharma, **Yiqiao Jin**, Vineeth Rakesh, Yingtong Dou, Menghai Pan, Mahashweta Das, and Srijan Kumar. *Sysformer: Safeguarding Frozen Large Language Models with Adaptive System Prompts.* In Proceedings of ICLR2026.
Acceptance rate: 28%.

6. Guancheng Wan*, Lucheng Fu*, Haoxin Liu, **Yiqiao Jin**, Hejia Geng, Eric Hanchen Jiang, Hui Yi Leong, Jinhe Bi, Yunpu Ma, Xiangru Tang, B. Aditya Prakash, Yizhou Sun, Wei Wang. *Beyond Magic Words: Sharpness-Aware Prompt Evolving for Robust Large Language Models.* In Proceedings of ICLR2026.
Acceptance rate: 28%.

7. Vibhor Agarwal, **Yiqiao Jin**, Mohit Chandra, Munmun De Choudhury, Srijan Kumar, Nishanth Sastry. *MedHalu: Hallucinations in Responses to Healthcare Queries by Large Language Models.*
In Proceedings of ICWSM2026.

8. Yijia Xiao, Wanjia Zhao, Junkai Zhang, **Yiqiao Jin**, Han Zhang, Zhicheng Ren, Renliang Sun, Haixin Wang, Guancheng Wan, Pan Lu, Xiao Luo, Yu Zhang, James Zou, Yizhou Sun, Wei Wang. *Protein Large Language Models: A Comprehensive Survey*. In Proceedings of EMNLP2025.
Acceptance rate: 22.2%.

9. Junyu Luo, Bohan Wu, Xiao Luo, Zhiping Xiao, **Yiqiao Jin**, Rong-Cheng Tu, Nan Yin, Yifan Wang, Jingyang Yuan, Wei Ju, Ming Zhang. *A Survey on Efficient LLM Training: From Data-centric Perspectives.* In Proceedings of ACL2025 main conference.

10. Qinlin Zhao, Jindong Wang, Yixuan Zhang, **Yiqiao Jin**, Kaijie Zhu, Hao Chen, Xing Xie. *CompeteAI: Understanding the Competition Behaviors in Large Language Model-based Agents.* In Proceedings of ICML2024, **Oral Presentation**.
Acceptance rate: 27.5%.

11. Yijia Xiao, **Yiqiao Jin**, Yushi Bai, Yue Wu, Xianjun Yang, Xiao Luo, Wenchao Yu, Xujiang Zhao, Yanchi Liu, Haifeng Chen, Wei Wang, Wei Cheng. *PrivacyMind: Large Language Models Can Be Contextual Privacy Protection Learners.* In Proceedings of EMNLP2024 main track and SoCal NLP Symposium 2024.
Acceptance rate: 20.8%.

12. Jinghan Zhang, Xiting Wang, **Yiqiao Jin**, Changyu Chen, Xinhao Zhang, Kunpeng Liu. *Prototypical Reward Network for Data-Efficient RLHF.* In Proceedings of ACL2024. **Oral Presentation**.
Acceptance rate: 23.5%.

13. Neng Kai Nigel Neo*, Yeon-Chang Lee*, **Yiqiao Jin**, Sang-Wook Kim, Srijan Kumar. *Towards Fair Graph Anomaly Detection: Problem, New Datasets, and Evaluation.* In Proceedings of CIKM2024.
Acceptance rate: 23.2%.

14. Changyu Chen, Xiting Wang, **Yiqiao Jin**, Victor Ye Dong, Li Dong, Rui Yan, Jim Cao, Yi Liu. *Semi-Offline Reinforcement Learning for Optimized Text Generation.* In Proceedings of ICML2023.
Acceptance rate: 27.9%.

15. Ruichao Yang, Xiting Wang, **Yiqiao Jin**, Chaozhuo Li, Jianxun Lian, Xing Xie. *Reinforcement Subgraph Reasoning for Fake News Detection.* In Proceedings of KDD2022.
Acceptance rate: 14.9%.

Journals

1. Chengyuan Deng, Yiqun Duan, Xin Jin, Heng Chang, Yijun Tian, Han Liu, Henry Peng Zou, **Yiqiao Jin**, Yijia Xiao, Yichen Wang, Shenghao Wu, Zongxing Xie, Kuofeng Gao, Sihong He, Jun Zhuang, Lu Cheng, Haohan Wang. *Deconstructing The Ethics of Large Language Models from Long-standing Issues to New-emerging Dilemmas.* Accepted at AI and Ethics.

Workshops

1. **Yiqiao Jin***, Yiyang Wang*, Lucheng Fu, Yijia Xiao, Yinyi Luo, Haoxin Liu, B. Aditya Prakash, Josiah Hester, Jindong Wang, and Srijan Kumar. Accepted at COLM 2026 LLA Workshop.

2. Lucheng Fu, Ye Yu, Yiyang Wang, **Yiqiao Jin**, Haibo Jin, B. Aditya Prakash, and Haohan Wang. *TextReg: Mitigating Prompt Distributional Overfitting via Regularized Text-Space Optimization*. Accepted at COLM 2026 Efficient Reasoning Workshop.

3. **Yiqiao Jin**, Yijia Xiao, Yiyang Wang, and Jindong Wang. *SciEvo: A 2 Million, 30-Year Cross-disciplinary Dataset for Temporal Scientometric Analysis.* Best Paper Award (2 out of 25 accepted papers) at AAAI2025 Good-Data Workshop.

4. **Yiqiao Jin**, Andrew Zhao, Yeon-Chang Lee, Meng Ye, Ajay Divakaran, and Srijan Kumar. *Empowering Interdisciplinary Insights with Dynamic Graph Embedding Trajectories.* Accepted at KDD2025 TGL Workshop.

5. **Yiqiao Jin**, Gang Wu, Yu Shen, and Stefano Petrangeli. *ScreenLLM: Stateful Screen Schema for Efficient Action Understanding and Prediction.* Accepted at WebConf2025 MM4SG Workshop.

6. Yijia Xiao, Edward Sun, **Yiqiao Jin**, Qifan Wang, and Wei Wang. *ProteinGPT: Multimodal LLM for Protein Property Prediction and Structure Understanding.* Accepted at ICML2025 FM4LS Workshop.

7. Sejoon Oh*, **Yiqiao Jin***, Megha Sharma, Donghyun Kim, Gaurav Verma, Eric Ma, Srijan Kumar. *UniGuard: Towards Universal Safety Guardrails for Jailbreak Attacks on Multimodal Large Language Models.* Accepted at AAAI2025 DAI Workshop.

8. Yijia Xiao, Edward Sun, **Yiqiao Jin**, Wei Wang. *RNA-GPT: Multimodal Generative System for RNA Sequence Understanding.* Accepted at NeurIPS2024 MLSB Workshop.

Preprints

1. Yinyi Luo, **Yiqiao Jin**, Weichen Yu, Mengqi Zhang, Srijan Kumar, Xiaoxiao Li, Weijie Xu, Xin Chen, Jindong Wang. *AgentArk: Distilling Multi-Agent Intelligence into a Single LLM Agent*. Under Review at NeurIPS2026.

2. Zhaolong Su*, Yinyi Luo*, **Yiqiao Jin***, Mengqi Zhang, Wenyue Hua, Srijan Kumar, Qingsong Wen, Jindong Wang. *Consistency Should Be the Priority for Unified Multimodal Models*. Under Review at ICML2026.

3. Jiaxi Yang*, Mengqi Zhang*, **Yiqiao Jin***, Hao Chen, Qingsong Wen, Lu Lin, Yi He, Srijan Kumar, Weijie Xu, James Evans, and Jindong Wang. *Topological Structure Learning Should Be A Research Priority for LLM-Based Multi-Agent Systems*. Under Review at JCST.

4. Kartik Sharma, **Yiqiao Jin**, Rakshit Trivedi, Srijan Kumar. *Efficient Knowledge Probing of Large Language Models by Adapting Pre-trained Embeddings.*

5. Yibo Hu, **Yiqiao Jin**, Meng Ye, Ajay Divakaran, Srijan Kumar. *The Influence of Text Variation on User Engagement in Cross-Platform Content Sharing.*

6. Junyu Luo, Weizhi Zhang, Ye Yuan, Yusheng Zhao, Junwei Yang, Yiyang Gu, Bohan Wu, Binqi Chen, Ziyue Qiao, Qingqing Long, Rongcheng Tu, Xiao Luo, Wei Ju, Zhiping Xiao, Yifan Wang, Meng Xiao, Chenwu Liu, Jingyang Yuan, Shichang Zhang, **Yiqiao Jin**, Fan Zhang, Xian Wu, Hanqing Zhao, Dacheng Tao, Philip S. Yu, and Ming Zhang. *Large Language Model Agent: A Survey on Methodology, Applications and Challenges*.

7. Ye Yuan, Junyu Luo, Guancheng Wan, Jinsheng Huang, Chengwu Liu, Junwei Yang, Yifang Qin, Zhiping Xiao, QingqingLong, Meng Xiao, **Yiqiao Jin**, Jianhong Tu, Yuqi Chen, Wei Ju, Zhongwei Wan, Yusheng Zhao, Xiao Luo, Yiwei Fu, Yizhou Sun, Wei Wang, Chenguang Wang, and Ming Zhang. *A Comprehensive Survey of Evaluating Multimodal Foundation Models: Hierarchical Perspective and Extensive Applications*. Under Review at ACL ARR.

Professional Experience

Amazon, Fulfillment By Amazon (FBA) June 2020 – Sep. 2020
Software Engineer Intern Seattle, WA

- Designed and implemented `IAR Manual Analysis`, a scalable and efficient workflow using `AWS Step Functions` and `AWS Lambda`. This service automates the aggregation of data points from multiple sources like `Amazon S3` and `DynamoDB` for `SageMaker` ML model training, handling $\geq$ 16,000 requests per summary stage;
- Automated the deployment of the workflow across all AWS Realms (EU/FE/NA) through `CloudFormation`;
- Establish `DataCraft` pipeline to enable automatic data ingestion from `DynamoDB` into the `Andes` dataset catalog, promoting broader internal adoption of these datasets for cross-functional teams and enhancing data accessibility;
- Perform `ablation analysis` on the inventory reconciliation model, identifying key performance bottlenecks and optimizing model performance.

IBM, China Development Laboratories June 2019 – Sep. 2019
Software Engineer Intern Beijing, China

- Developed `Compass DataRouter`, a routing service for `Compass` project based on `Golang` and `MongoDB`, reducing memory usage and accelerating data retrieval;
- Enhanced the monitoring dashboard for `Compass` towards a more intuitive and responsive user interface with `React.js`.

Services

Georgia Tech School of Computer Science Graduate Student Association (SCSGSA) Aug. 2024 –
Communications Chair Atlanta, GA

- Co-organize the College of Computing (CoC) Graduate Welcome Event with over 1,000 attendees, including new students, faculty, and alumni;
- Lead the design and maintenance of the SCSGSA website (`https://scsgsa.cc.gatech.edu/`);
- Promote SCSGSA events, including workshops, networking sessions, and professional development programs;

Academic Services

- Area Chair: ICLR2024/2023 Tiny Paper Track
- PC Member/Reviewer
  - **Conferences**:
    * AAAI 2026/2025/2024/2023;
    * ACL ARR (Feb. 2024 until Now, including ACL/EMNLP/NAACL);
    * AISTATS 2025/2024;
    * ASONAM 2024/2023;
    * CIKM 2024/2023;
    * COLM 2026/2025;
    * CVPR 2026/2025;
    * ICCV 2025;
    * ICLR 2026/2025/2024, ICLR 2024/2023 Tiny Paper;

    * ICML 2026/2025/2024;
    * KDD 2025/2024/2023;
    * LoG 2024/2023/2022;
    * NeurIPS 2026/2025/2024/2023, NeurIPS 2025/2024/2023 Datasets & Benchmarks Track;
    * SDM 2024;
    * Web Conference 2025/2024.
  - **Journals**:
    * ACM Transactions on Information Systems (TOIS); 2023 –
    * ACM Transactions on Knowledge Discovery from Data (TKDD); 2023 –
    * IEEE Transactions on Knowledge and Data Engineering (TKDE); 2022 –
    * ACM Transactions on Recommender Systems (TORS); 2022 –
    * ACM Transactions on Intelligent Systems and Technology (TIST); 2022 –
    * ACM Transactions on Social Computing (TSC); 2022 –
    * IEEE Transactions on Artificial Intelligence (TAI); 2023 –
    * PLOS ONE; 2023 –
    * International Journal of Data Science and Analytics (JDSA); 2021 –
    * SCIENCE CHINA Information Sciences (SCIS); 2022 –
    * The Journal of Data-centric Machine Learning Research (DMLR). 2023 –
  - **Workshops**
    * AFT@NeurIPS2023, Ai4Science@NeurIPS2023, Ai4D3@NeurIPS2023, GenBio@NeurIPS2023, TGL Workshop@NeurIPS2023;
    * SPIGM@ICML2023.

- Volunteer: NAACL2025, AAAI2025, SC2024, VLDB2024

Honors and Awards

- MLCommons ML and Systems Rising Stars. 2026
- Best Paper Award at Good-Data @ AAAI2025 Workshop. 2025
- Roblox Graduate Fellowship Finalist. 2024
- AAAI Student Scholarship (3 times). 2025, 2023, 2022
- Microsoft Research “Star of Tomorrow” Award of Excellence. 2021
- UCLA Dean’s Honor List for Superior Academic Achievement (5 times). 2019 – 2021
  - Spring 2019, Winter 2020, Spring 2020, Winter 2021, Spring 2021