Yiqiao Jin (Ahren)

Last Updated: Sep. 22nd, 2025

CONTACT

Georgia Institute of Technology
Atlanta, GA 30332
WeChat: go_bruins
Twitter: AhrenJin
github.com/Ahren09
(+86)189-1119-7743
(470)962-0241

ahren09.github.io
linkedin.com/in/ahren-jin
scholar.google.com/citations?user=eY85qm4AAAAJ
Semantic Scholar: 2087723977
ORCID: 0000-0002-6974-5970
yjin328@gatech.edu
ahren2040@g.ucla.edu

RESEARCH INTERESTS

- Large Language Models, LLM Agents, Multimodal Learning, Graph Neural Networks, Computational Social Science.

EDUCATION

Georgia Institute of Technology Aug. 2022 – Present

- Ph.D. Candidate, Computer Science. GPA: **4.0/4.0**.

University of California, Los Angeles (UCLA) Sep. 2018 – Dec. 2021

- B.S., Computer Science. GPA: **3.82/4.0**.
- Published 4 papers (including 3 first-authored papers) at top-tier machine learning and data mining conferences (AAAI, KDD, Web Conference).

RESEARCH EXPERIENCE

Georgia Tech College of Computing, School of CSE Aug 2022 – Present
Graduate Research Assistant Atlanta, GA
Advisor: Dr. Srijan Kumar

- Research Topics: Large Language Models (EMNLP'25, WebConf'24, ACL'24), LLM Robustness and Safety (Under Review at AAAI'26), Multimodal LLMs (CVPR'25, ACL'24), Recommender Systems and Dynamic Graph Mining (KDD'23), Social Network Analysis (CIKM'24, KDD'23), Fair Graph Mining (CIKM'24).

J.P. Morgan AI Research May 2025 – Aug 2025
Research Scientist Intern New York, NY
Mentors: Dr. Zhen Zeng, Dr. Sumitra Ganesh.

- Research Topics: Large Language Models for Infographics Understanding, GUI Agent.
- Co-authored paper is under review at ACL ARR.

Visa Research Aug 2024 – May 2025
Research Collaborator Remote
Mentors: Dr. Vineeth Rakesh Mohan, Dr. Yingtong Dou, Dr. Menghai Pan.

- Research Topics: Large Language Models, Retrieval-augmented Generation, Context Compression, Efficient NLP.
- Co-authored paper is under review at ACL ARR.

Adobe Research May 2024 – Aug 2024
Research Scientist Intern San Jose, CA
Advisors: Dr. Gang Wu, Dr. Yu Shen, Dr. Stefano Petrangeli

- Research Topics: Multimodal Large Language Models (MLLMs) Fine-tuning, GUI Agents, GUI and Video Tutorial Understanding (WebConf'25).

Microsoft Research Asia (MSRA), Social Computing Group Dec. 2020 – July 2022
Research Scientist Intern Beijing, China
Advisors: Dr. Xiting Wang, Dr. Jindong Wang, and Dr. Xing Xie

- Research Topics: Large Language Models (EMNLP'24, ICML'24, ICML'23, AAAI'23), LLM Agents (EMNLP'24, ICML'24), Multicultural and Multimodal LLM, Scientometric Analysis, Computational Social Science, Misinformation Detection (KDD'22, AAAI'22), Few-shot Learning (ACL'24, AAAI'23) , Explainable AI (AAAI'22).

UCLA, Scalable Analytics Institute (ScAi) June 2021 – June 2022
Undergraduate Research Assistant Los Angeles, CA
Advisors: Dr. Yizhou Sun and Dr. Wei Wang

- Research Topics: Large Language Models (EMNLP'25, ACL'25, EMNLP'24), Graph Neural Networks and Data Mining (WWW'23), LLM Fine-tuning, Recommender Systems (WWW'23).
- I am continuously collaborating with UCLA researchers during my Ph.D. studies.

Publications

Conference Papers

1. **Yiqiao Jin***, Qinlin Zhao*, Yiyang Wang, Hao Chen, Kaijie Zhu, Yijia Xiao, Jindong Wang. *AgentReview: Exploring Peer Review Dynamics with LLM Agents.* In Proceedings of EMNLP'24 main track **Oral Presentation** and SoCal NLP Symposium 2024.
   Acceptance rate: 20.8%.

2. **Yiqiao Jin**, Minje Choi, Gaurav Verma, Jindong Wang, Srijan Kumar. *MM-Soc: Benchmarking Multimodal Large Language Models in Social Media Platforms.* In Proceedings of ACL'24 and the 3rd ALVR Workshop.
   Acceptance rate: 23.5%.

3. **Yiqiao Jin***, Mohit Chandra*, Gaurav Verma, Yibo Hu, Munmun De Choudhury, Srijan Kumar. *Better to Ask in English: Cross-Lingual Evaluation of Large Language Models for Healthcare Queries.* In Proceedings of the ACM Web Conference 2024. **Oral Presentation**.
   Acceptance rate: 20.2%.
   Media Coverage by Scientific American, The World, and Georgia Tech.

4. **Yiqiao Jin**, Yeon-Chang Lee, Kartik Sharma, Meng Ye, Karan Sikka, Ajay Divakaran, Srijan Kumar. *Predicting Information Pathways Across Online Communities.* In Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD'23). **Oral Presentation**.
   Acceptance rate: 22.1%.

5. **Yiqiao Jin**, Yunsheng Bai, Yanqiao Zhu, Yizhou Sun, Wei Wang. *Code Recommendation for Open Source Project Developers.* In Proceedings of the ACM Web Conference 2023. **Oral Presentation**.
   Acceptance rate: 19.2%

6. **Yiqiao Jin**, Xiting Wang, Yaru Hao, Yizhou Sun, Xing Xie. *Prototypical Fine-tuning: Towards Robust Performance Under Varying Data Sizes.* In Proceedings of the 37th AAAI Conference (AAAI'23). **Oral Presentation**.
   Acceptance rate: 19.6%.

7. **Yiqiao Jin**, Xiting Wang, Ruichao Yang, Yizhou Sun, Wei Wang, Hao Liao, Xing Xie. *Towards Fine-Grained Reasoning for Fake News Detection.* In Proceedings of the 36th AAAI Conference (AAAI'22). **Oral Presentation**.
   Acceptance rate: 14.6%.

8. Junyu Luo, Bohan Wu, Xiao Luo, Zhiping Xiao, **Yiqiao Jin**, Rong-Cheng Tu, Nan Yin, Yifan Wang, Jingyang Yuan, Wei Ju, Ming Zhang. *A Survey on Efficient LLM Training: From Data-centric Perspectives.* In Proceedings of ACL'25.

9. Qinlin Zhao, Jindong Wang, Yixuan Zhang, **Yiqiao Jin**, Kaijie Zhu, Hao Chen, Xing Xie. *CompeteAI: Understanding the Competition Behaviors in Large Language Model-based Agents.* In Proceedings of the 41st International Conference on Machine Learning (ICML'24), **Oral Presentation**.
   Acceptance rate: 27.5%.

10. Yijia Xiao, **Yiqiao Jin**, Yushi Bai, Yue Wu, Xianjun Yang, Xiao Luo, Wenchao Yu, Xujiang Zhao, Yanchi Liu, Haifeng Chen, Wei Wang, Wei Cheng. *Large Language Models Can Be Good Privacy Protection Learners.* In Proceedings of EMNLP'24 main track and SoCal NLP Symposium 2024.
    Acceptance rate: 20.8%.

11. Jinghan Zhang, Xiting Wang, **Yiqiao Jin**, Changyu Chen, Xinhao Zhang, Kunpeng Liu. *Prototypical Reward Network for Data-Efficient RLHF.* In Proceedings of ACL'24. **Oral Presentation**.
    Acceptance rate: 23.5%.

12. Vibhor Agarwal, **Yiqiao Jin**, Mohit Chandra, Munmun De Choudhury, Srijan Kumar, Nishanth Sastry *MedHalu: Hallucinations in Responses to Healthcare Queries by Large Language Models.* In Proceedings of ICWSM'26.

13. Yijia Xiao, Wanjia Zhao, Junkai Zhang, **Yiqiao Jin**, Han Zhang, Zhicheng Ren, Renliang Sun, Haixin Wang, Guancheng Wan, Pan Lu, Xiao Luo, Yu Zhang, James Zou, Yizhou Sun, Wei Wang. *Protein Large Language Models: A Comprehensive Survey*. In Proceedings of EMNLP'25.
    Acceptance rate: 22.2%.

14. Neng Kai Nigel Neo*, Yeon-Chang Lee*, **Yiqiao Jin**, Sang-Wook Kim, Srijan Kumar. *Towards Fair Graph Anomaly Detection: Problem, New Datasets, and Evaluation.* In Proceedings of the 33rd ACM International Conference on Information and Knowledge Management (CIKM'24). Acceptance rate: 23.2%.

15. Changyu Chen, Xiting Wang, **Yiqiao Jin**, Victor Ye Dong, Li Dong, Rui Yan, Jim Cao, Yi Liu. *Semi-Offline Reinforcement Learning for Optimized Text Generation.* In Proceedings of the 40th International Conference on Machine Learning (ICML'23). Acceptance rate: 27.9%.

16. Ruichao Yang, Xiting Wang, **Yiqiao Jin**, Chaozhuo Li, Jianxun Lian, Xing Xie. *Reinforcement Subgraph Reasoning for Fake News Detection.* In Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD'22). Acceptance rate: 14.9%.

Journals

1. Chengyuan Deng, Yiqun Duan, Xin Jin, Heng Chang, Yijun Tian, Han Liu, Henry Peng Zou, **Yiqiao Jin**, Yijia Xiao, Yichen Wang, Shenghao Wu, Zongxing Xie, Kuofeng Gao, Sihong He, Jun Zhuang, Lu Cheng, Haohan Wang. *Deconstructing The Ethics of Large Language Models from Long-standing Issues to New-emerging Dilemmas.* Accepted at AI and Ethics.

Workshops

1. **Yiqiao Jin**, Yijia Xiao, Yiyang Wang, Jindong Wang. *SciEvo: A 2 Million, 30-Year Cross-disciplinary Dataset for Temporal Scientometric Analysis.* Best Paper Award (2 out of 25 accepted papers) at AAAI'25 Good-Data Workshop.

2. **Yiqiao Jin**, Kartik Sharma, Vineeth Rakesh, Yingtong Dou, Menghai Pan, Mahashweta Das, and Srijan Kumar. *SARA: Selective and Adaptive Retrieval-augmented Generation with Context Compression.* Accepted at NeurIPS'25 ES-FoMo III Workshop.

3. **Yiqiao Jin**, Andrew Zhao, Yeon-Chang Lee, Meng Ye, Ajay Divakaran, Srijan Kumar. *Empowering Interdisciplinary Insights with Dynamic Graph Embedding Trajectories.* Accepted at KDD'25 TGL Workshop.

4. **Yiqiao Jin**, Gang Wu, Yu Shen, and Stefano Petrangeli, *Stateful Screen Schema for Efficient Action Understanding and Prediction.* Accepted at WebConf'25 MM4SG Workshop.

5. Kartik Sharma, **Yiqiao Jin**, Vineeth Rakesh, Yingtong Dou, Menghai Pan, Mahashweta Das, and Srijan Kumar. *Sysformer: Safeguarding Frozen Large Language Models with Adaptive System Prompts.* Accepted at ICML'25 R2-FM Workshop.

6. Yijia Xiao, Edward Sun, **Yiqiao Jin**, Qifan Wang, Wei Wang. *ProteinGPT: Multimodal LLM for Protein Property Prediction and Structure Understanding.* Accepted at ICML'25 FM4LS Workshop.

7. Shudong Liu, **Yiqiao Jin**, Cheng Li, Derek F. Wong, Qingsong Wen, Lichao Sun, Haipeng Chen, Xing Xie, Jindong Wang. *CultureVLM: Characterizing and Improving Cultural Understanding of Vision-Language Models for over 100 Countries.* Accepted at CVPR'25 VLMs4All Workshop.

8. Sejoon Oh*, **Yiqiao Jin***, Megha Sharma, Donghyun Kim, Gaurav Verma, Eric Ma, Srijan Kumar. *UniGuard: Multimodal Safety Guardrails for Jailbreak Attacks on Multimodal Large Language Models.* Accepted at AAAI'25 DAI Workshop.

9. Yijia Xiao, Edward Sun, **Yiqiao Jin**, Wei Wang. *RNA-GPT: Multimodal Generative System for RNA Sequence Understanding.* Accepted at NeurIPS'24 MLSB Workshop.

Preprints

1. Jiaxi Yang*, Mengqi Zhang*, Yiqiao Jin*, Hao Chen, Qingsong Wen, Lu Lin, Yi He, Weijie Xu, James Evans, Jindong Wang. *Topological Structure Learning Should Be A Research Priority for LLM-Based Multi-Agent Systems*. Under Review at NeurIPS'25.

2. Mohit Chandra, Siddharth Sriraman, Harneet Singh Khanuja, Yiqiao Jin, Munmun De Choudhury. *Reasoning Is Not All You Need: Examining LLMs for Multi-Turn Mental Health Conversations*. Under Review at ACL ARR.

3. Kartik Sharma, **Yiqiao Jin**, Rakshit Trivedi, Srijan Kumar. *Efficient Knowledge Probing of Large Language Models by Adapting Pre-trained Embeddings.* Under Review at AAAI'26.

4. Guancheng Wan, Lucheng Fu, Haoxin Liu, **Yiqiao Jin**, Hejia Geng, Eric Hanchen Jiang, Hui Yi Leong, Jinhe Bi, Yunpu Ma, Xiangru Tang, B. Aditya Prakash, Yizhou Sun, Wei Wang. *Beyond Magic Words: Sharpness-Aware Prompt Evolving for Robust Large Language Models.* Under Review at ICLR'26.

5. Varun Vankineni, Sai Prasath Suresh, Gaurav Verma, **Yiqiao Jin**, Xiaojing Wang, Rajarajeswari Balasubramaniyan, Srijan Kumar. *Investigating the Effect of Automated vs. Human-Verified Adversarial Samples on Improving the Robustness of LLMs.*

Professional Experience

Amazon, Fulfillment By Amazon (FBA) June 2020 – Sep. 2020
Software Engineer Intern Seattle, WA

- Designed and implemented `IAR Manual Analysis`, a scalable and efficient workflow using `AWS Step Functions` and `AWS Lambda`. This service automates the aggregation of data points from multiple sources like `Amazon S3` and `DynamoDB` for `SageMaker` ML model training, handling $\geq$ 16,000 requests per summary stage;
- Automated the deployment of the workflow across all AWS Realms (EU/FE/NA) through `CloudFormation`;
- Establish `DataCraft` pipeline to enable automatic data ingestion from `DynamoDB` into the `Andes` dataset catalog, promoting broader internal adoption of these datasets for cross-functional teams and enhancing data accessibility;
- Perform `ablation analysis` on the inventory reconciliation model, identifying key performance bottlenecks and optimizing model performance.

IBM, China Development Laboratories June 2019 – Sep. 2019
Software Engineer Intern Beijing, China

- Developed `Compass DataRouter`, a routing service for `Compass` project based on `Golang` and `MongoDB`, reducing memory usage and accelerating data retrieval;
- Enhanced the monitoring dashboard for `Compass` towards a more intuitive and responsive user interface with `React.js`.

Services

Georgia Tech School of Computer Science Graduate Student Association (SCSGSA) Aug. 2024 –
Communications Chair Atlanta, GA

- Co-organize the College of Computing (CoC) Graduate Welcome Event with over 1,000 attendees, including new students, faculty, and alumni;
- Lead the design and maintenance of the SCSGSA website (`https://scsgsa.cc.gatech.edu/`);
- Promote SCSGSA events, including workshops, networking sessions, and professional development programs;

Academic Services

- Area Chair: ICLR'24/23 Tiny Paper Track
- PC Member/Reviewer
  - **Conferences**:
    - AAAI'26/25/24/23;
    - ACL ARR (Including ACL/EMNLP/NAACL 26/25/24);
    - AISTATS'25/24;
    - ASONAM'24/23;
    - CIKM'24/23;
    - COLM'25;
    - CVPR'25;
    - ICCV'25;
    - ICLR'25/24, ICLR'24/23 Tiny Paper;
    - ICML'25/24;
    - KDD'25/24/23;
    - LoG'24/23/22;
    - NeurIPS'25/24/23, NeurIPS'25/24/23 Datasets & Benchmarks Track;
    - SDM'24;

    - Web Conference'24.
  - **Journals**:
    - ACM Transactions on Information Systems (TOIS); 2023 –
    - ACM Transactions on Knowledge Discovery from Data (TKDD); 2023 –
    - IEEE Transactions on Knowledge and Data Engineering (TKDE); 2022 –
    - ACM Transactions on Recommender Systems (TORS); 2022 –
    - ACM Transactions on Intelligent Systems and Technology (TIST); 2022 –
    - ACM Transactions on Social Computing (TSC); 2022 –
    - IEEE Transactions on Artificial Intelligence (TAI); 2023 –
    - PLOS ONE; 2023 –
    - International Journal of Data Science and Analytics (JDSA); 2021 –
    - SCIENCE CHINA Information Sciences (SCIS); 2022 –
    - The Journal of Data-centric Machine Learning Research (DMLR). 2023 –
  - **Workshops**
    - AFT@NeurIPS2023, Ai4Science@NeurIPS2023, Ai4D3@NeurIPS2023, GenBio@NeurIPS2023, TGL Workshop@NeurIPS2023;
    - SPIGM@ICML2023.
- Volunteer: NAACL'25, AAAI'25, VLDB'24

Honors and Awards

- Best Paper Award at Good-Data @ AAAI'25 Workshop. 2025
- Roblox Graduate Fellowship Finalist. 2024
- AAAI Student Scholarship (3 times). 2025, 2023, 2022
- Microsoft Research "Star of Tomorrow" Award of Excellence. 2021
- UCLA Dean's Honor List for Superior Academic Achievement (5 times). 2019 – 2021
  - Spring 2019, Winter 2020, Spring 2020, Winter 2021, Spring 2021